



HANDVANTAGE

A REFERENCE FOR PROCUREMENT COMMITTEES

The Agentic AI Procurement *Handbook.*

Sector dossiers, persona briefs, and the architectural reference for buying and deploying agentic AI in regulated enterprises.

A peer-to-peer reference, not a sales document.

This handbook compiles the editorial published at **workspace.handvantage.com/insights** into a single reference for procurement committees evaluating agentic AI for deployment in regulated environments. It is written for senior buyers — CISOs, CFOs, COOs, compliance officers, and the founders and boards they brief.

The content is human-authored under the founder's byline. Every regulatory citation references published guidance — by article number, by clause number, by document number — so a reader who wants to verify any claim has the path to do so.

The handbook is published openly. No email gate. No registration. Reference it, share it, attach it to your committee deck, cite it in your own writing. The credit and the link back to the source are appreciated; the use is not restricted.

PUBLISHED: 2026-05-11

VERSION: 1.0

AUTHOR: JOSH OLAYEMI, FOUNDER, HANDVANTAGE

CONTACT: [HELLO@HANDVANTAGE.COM](mailto:hello@handvantage.com)

CANONICAL URL: [WORKSPACE.HANDVANTAGE.COM](https://workspace.handvantage.com)

LICENCE: FREE TO SHARE, CITE, AND FORWARD

Table of *contents*.

PART I

The procurement landscape

01	Why agentic AI procurement keeps stalling	1
02	The five-voice procurement negotiation	—

PART II

For specific roles

03	For the CISO	—
04	For the CFO	—
05	For the CEO and the board	—

PART III

Sector dossiers

06	Financial services — FINRA, SEC, the audit-trail problem	—
07	Healthcare — HIPAA, FDA SaMD, the supervision question	—
08	Fintech — BSA/AML, fair lending, the sponsor-bank question	—
09	Canadian public sector — TBS Directive, sovereignty, Indigenous data	—
10	Legal services — privilege, competence, the supervision rule	—

PART IV

Architecture and posture

I

The procurement landscape.

Why so many agentic AI deployments stall at the compliance review, and the five-voice negotiation underneath every procurement decision in a regulated enterprise.

PART I

CHAPTER 01

Why agentic AI procurement *keeps* *stalling*.

Roughly 40% of agentic AI projects are being cancelled or shelved, often after substantial investment. The pattern in cancelled projects is consistent — and it isn't about technology fit.

Gartner's research recently put a number on something we'd been seeing in the field: roughly 40% of agentic AI projects are being cancelled or shelved, often after substantial investment. The headline reaction is to point at technology immaturity. The actual pattern, when we look at the cancelled projects we have visibility into, is consistently about governance — not about whether the technology works, but about whether the organisation can defend its use of the technology.

Cancelled projects fall into three rough categories. The first is the technology-fit cancellation: the platform did not deliver what was promised on the demo. These are real, but rarer than the headline suggests. The second is the cost-fit cancellation: the platform delivered, but the unit economics didn't work at production scale. The third — and most common in our sample — is the governance-fit cancellation: the platform delivered, the costs penciled, but the security or compliance team could not approve a production rollout because the evidence story didn't exist.

The third — and most common in our sample — is the governance-fit cancellation: the platform delivered, the costs penciled, but the security or compliance team could not approve a production rollout because the evidence story didn't exist.

— WHY 40% OF AGENTIC AI PROJECTS FAIL (AND THE GOVERNANCE ANSWER)

GARTNER FORECAST • 2027

40%

of agentic AI projects cancelled, almost always for missing evidence, not missing controls.

The governance-fit cancellation has a recurring shape. The technology team builds a proof of concept, demonstrates value, and submits a production-rollout proposal. The compliance team reviews the proposal and asks for evidence — not for whether the platform can do prompt-injection defence in a demo, but for whether the platform produces contemporaneous, machine-verifiable, framework-mapped evidence of every action it takes. The technology team turns to the vendor. The vendor says some version of "we generate logs, you can build a compliance pipeline on top." The compliance team responds: "we don't have the budget, the people, or the time to build a compliance pipeline; the platform should produce the evidence directly." The proposal stalls. After two or three revisions, the proposal is shelved.

The shape is consistent enough that we treat it as the default failure mode for any agentic AI project where the buying organisation is in a regulated sector. The pattern is not a technology problem. The platforms work. The pattern is a structural mismatch between what vendors are selling (a product layer) and what regulated buyers need (a product layer plus an evidence layer plus a control-mapping layer plus a verification layer).

The governance answer is to build the evidence layer into the platform from the architectural premise, not as an add-on. This is what the 7-Layer Defence Architecture is — every layer produces audit events, every event is signed, every event is mapped to specific framework controls, every assessment is run against the runtime evidence rather than against an attestation document. The architecture is structural; the evidence is contemporaneous; the assessment is continuous.

This shape doesn't make the project's compliance review trivial. It makes the review possible. The conversation between the technology team and the compliance team becomes "here is the Trust Report scoped to your frameworks, with your audit window" rather than "we'll figure out how to prove this once we're in production." The proposal moves forward. The project ships. The 40% failure rate moves the other way.

We are not arguing that this is the only design pattern that works. We are arguing it is the design pattern that addresses the actual failure mode in cancelled agentic AI projects. The vendors who continue to ship a product layer alone will continue to lose deals to vendors who ship the evidence layer alongside it. The buyers who continue to evaluate platforms on demo behaviour alone will continue to cancel projects when the compliance review reaches the evidence question. The pattern is settled enough now that we expect it to continue.

PART I

CHAPTER 02

Five people in the room. *Five reads of the same platform.*

Most agentic AI procurement decisions stall not because anyone in the room is wrong, but because the people in the room are answering five different questions — and rarely realise it.

Most agentic AI procurement decisions stall not because anyone in the room is wrong, but because the people in the room are answering five different questions. The CEO is asking whether the strategy is defensible to a board. The CFO is asking whether the line item consolidates. The COO is asking what changes for throughput. The CISO is asking whether the security and compliance posture holds. The VP Sales is asking whether the cycle gets faster. The department lead is asking what their team does on Monday morning.

All five answers are good ones. The platform delivers all five — but in different vocabularies, with different proofs, on different timelines. The translation gap is what kills deals. A vendor that pitches the CISO with the CFO's vocabulary loses the CISO's trust. A vendor that pitches the CFO with the CISO's vocabulary loses the CFO's attention. The same platform, when described in the language each persona actually uses, lands differently with each.

All five answers are good ones. The platform delivers all five — but in different vocabularies, with different proofs, on different timelines.

— FOR THE PROCUREMENT COMMITTEE

THE FIVE VOICES

CEO • CFO • COO • CISO • VP Sales

Each is answering a different question. The deal moves when the language for all five lines up.

The CEO's read is strategic. The version of the sentence that holds: "We deployed a sovereign AI platform on infrastructure we control, graded continuously across eleven regulatory frameworks, with one identity for our team and one signed audit trail for the auditor." Every clause is something the platform actually delivers; none of it requires the CEO to apologise for vendor sprawl, missing controls, or absence of evidence. This matters because the alternative — "we're evaluating several AI assistants and hope to have a recommendation by Q4" — is increasingly read by boards as a governance failure rather than a capability decision.

The CFO's read is consolidation. Eight productivity vendor contracts become one MSA. Identity is included via Keycloak preconfigured (federate to Okta, Microsoft Entra ID, Auth0, or Google Workspace via OIDC if one's in place). The SIEM is the only piece that genuinely stays separate — the platform feeds your SIEM rather than replacing it, because specialised security observability is its own category. The contract surface shrinks; the security review surface shrinks; the renewal cycle shrinks. The math the CFO can verify against last year's vendor master list, not a vendor-marketing claim.

The COO's read is operational. The five named agents (Pilot, Hunter, Sentry, Concierge, Analyst) handle work that previously needed junior-staff time. The capacity model that matters: agents add throughput at the per-task level while humans stay accountable at the per-decision level. The honest caveat: agents are not a headcount substitute on the day of deployment. The capacity gain shows up over the second and third quarters of use, as workflows are tuned and the team learns which decisions to delegate. Vendors who promise a headcount cut on day one are selling something we do not.

The CISO's read is structural. Every prompt mediated, every action signed, every event mapped to controls under eleven frameworks. The audit log is non-contestable on whether the layers were running — they were, by construction. The platform produces the runtime evidence record; the customer's management system still owns the policy, procedures, and management review. The platform reduces the operational burden of compliance — it does not deliver compliance as a finished output.

The VP Sales's read is cycle time. The Hunter agent generates the customer brief from the workspace's knowledge surface. The Analyst cross-references the customer's tech stack. The Documents pillar drafts the proposal against the team's template, populated with the brief's facts. The Concierge assembles the procurement-friendly attachments from the Files pillar. The cycle from qualified inquiry to proposal-in-customer-hand compresses from three-to-seven business days to one-to-three. The rep's time allocation moves from approximately 30% selling and 70% paperwork to roughly 60% selling and 40% review.

When the five readings of the platform line up — the CEO has a defensible board narrative, the CFO has a single contract, the COO has a phased capacity plan, the CISO has the audit posture, the VP Sales has the cycle-time savings — the procurement decision moves from "maybe in Q4" to "let's start with the pilot deployment." This is the difference between platforms that ship and platforms that get cancelled.

The procurement reality this handbook addresses: each persona deserves a reading of the platform in their own vocabulary. The chapters that follow do that — Part II for the persona briefs, Part III for the sector dossiers, Part IV for the architectural and compliance posture that anchors all of them.

III

Sector dossiers.

One brief per regulated sector. Each names the operative regulatory frames, walks the deployment patterns, and clarifies what the platform fits and what it deliberately does not — with named alternative vendor categories for the use cases that aren't ours.

PART III

CHAPTER 06

Financial services — *FINRA, SEC, the audit-trail problem.*

Three regulatory regimes converge on one technical question: can the firm reconstruct, on demand, the exact decision sequence an AI agent followed when interacting with a client account? Most platforms cannot. The deals are stalling there.

Financial services is the sector where agentic AI buying decisions stall longest, and the reason is consistent across the firms we have visibility into. The procurement question is not whether the platform can perform the task. The procurement question is whether the firm can defend the platform's use to three separate regulators on three separate frameworks, with the audit trail an examiner will ask for in the first thirty minutes of any examination. Most platforms cannot produce that audit trail. This dossier is the regulatory landscape, the specific questions an examiner will ask, and the structural gap that is killing deals at the compliance review.

The three regulatory regimes a US-domiciled financial services firm has to satisfy simultaneously when deploying agentic AI: FINRA (for broker-dealers, member firms), the SEC (for registered investment advisers, investment companies, and the broker-dealer side via Reg BI and Rule 17a), and applicable state regulators (NY DFS, California DFPI, etc., where the firm is licensed). Each has its own framework. None of them have published agentic-AI-specific rules. All of them apply existing rules to AI deployments and let the firm work out the implementation. This is the consistent failure mode: the rules apply unambiguously; the implementation guidance is the firm's problem.

The platform's logs are necessary. They are not the supervisory record on their own. The firm's compliance team needs the platform to produce a unified, signed, retention-compliant interaction record per client.

— THE AUDIT-TRAIL PROBLEM

THE OPERATIVE RULES

FINRA 3110 · 4511 · SEC 17a-4 · Reg BI

Four federal rules apply to any AI interaction touching a client account. None were written for AI. All apply unmodified.

FINRA Rule 3110 (Supervision) is the operative rule for agentic AI in a member firm. The rule requires the firm to establish and maintain a supervisory system reasonably designed to achieve compliance with applicable securities laws. When an AI agent takes an action that would have been a registered representative action — drafting client communications, conducting account analysis, generating recommendations, interacting with order management — the supervision rule applies to that action. The supervisor of record has to have the means to review what the agent did, why it did it, and whether it complied with applicable rules. This is structural: the firm cannot delegate the supervisory function to the platform; the firm has to maintain it.

FINRA Rule 4511 (Books and Records) and SEC Rule 17a-4 (records retention) are the operative rules for the audit trail. Rule 4511 requires firms to maintain books and records as required by SEC Rules 17a-3 and 17a-4. Rule 17a-4 requires that broker-dealer records be preserved in a non-rewriteable, non-erasable format (the WORM standard, Write Once Read Many) for the prescribed retention periods. When the AI agent's decision sequence is part of the supervisory record, that record needs to meet the WORM standard and the retention period. Logs that are mutable, or that can be silently filtered, do not satisfy 17a-4.

SEC Reg BI (Regulation Best Interest, 17 CFR § 240.15l-1) requires that recommendations to retail customers be in the customer's best interest. When an AI agent generates a recommendation, the Reg BI obligations apply to the recommendation. The firm has to be able to demonstrate the basis for the recommendation, the conflicts of interest considered, the disclosures made — and the demonstration has to survive an SEC examination, which means the underlying record has to be reconstructable on demand from the firm's records.

FINRA Regulatory Notice 24-09 (June 2024, on supervisory considerations for the use of generative AI) and the 2025 Annual Regulatory Oversight Report's AI section establish FINRA's expectation that firms apply existing rules to AI deployments and document the supervisory framework. The notice does not introduce new rules; it makes explicit that the firm cannot avoid Rule 3110 or Rule 4511 obligations by claiming the AI is autonomous.

The pattern in recent FINRA examinations of firms with AI deployments has been consistent. First question: "Show me the AI inventory — every system, every business function, every supervisor of record." Second question: "For a sample client interaction, show me the full decision record — the prompt, the agent's reasoning trace, the tools the agent invoked, the data the agent retrieved, and the final action taken — with timestamps and the supervisor sign-off where required." This is where most firms fail. The platform produces some of these elements, the firm's logging captures some, the data warehouse retains some, but the firm cannot, on demand, produce the unified decision record. Reconstructing it from disconnected logs takes weeks. The examiner expects it within the examination period.

What the audit-trail problem looks like in practice. A broker-dealer deploys an agentic AI assistant for registered representatives. The assistant drafts client communications, analyses account positions, suggests rebalancing actions for the rep to consider. Each of these is potentially a Reg BI moment, a 3110 supervision moment, and a 4511/17a-4 record. The platform vendor produces a chat log. The platform vendor produces a tool-call log. The platform vendor produces a model-inference log. Three logs. None of them are joined at the platform level. The firm's SIEM ingests them but doesn't reassemble them into a per-interaction record. When the examiner asks for one client's six-month interaction history, the firm has to reconstruct it from three sources.

The structural answer is to treat the interaction record as a first-class output of the platform, signed at production, retained per the WORM standard, and exported to the firm's record-keeping infrastructure in a format the firm's compliance team can hand to an examiner without rework. This is what an audit-grade record looks like in practice: per-interaction, with the prompt, the agent's reasoning trace, the tool invocations, the data retrievals, the final action, the supervisor sign-off (where required), and the policy version active at the time. RFC 3161 timestamps. Merkle-tree anchoring (so the firm can prove the record was not altered after the fact, which is what 17a-4's non-rewriteable requirement is about).

What the firm still owns. The supervisory framework itself is the firm's. The platform produces the records; the firm's compliance team determines which interactions require supervisor review, what the review consists of, and how the review is documented. The Written Supervisory Procedures (WSPs) need to address AI explicitly, with the supervisor of record named for each AI system in the inventory. The firm's training programme has to cover AI use; the firm's annual compliance certification has to cover AI controls.

What firms ahead of this curve are doing. They have built the AI inventory before deploying the agentic system, not after. They have updated their WSPs to address AI before the system goes live, with review by their compliance counsel, not just internal compliance. They have demanded the unified interaction record from the platform vendor in the procurement process; they have not accepted "we produce logs" as an answer. The vendors who can demonstrate the record win the deal. The vendors who cannot are losing.

The state-regulator dimension. NY DFS Cybersecurity Regulation (23 NYCRR 500), as amended in 2023, requires written cybersecurity policies and risk assessments for AI systems. California Consumer Privacy Act and California Privacy Rights Act extend privacy obligations to automated decision-making systems. The audit trail that satisfies FINRA satisfies most of the state requirements as a side effect, because the state requirements largely require a documented, retained record of AI-driven decisions touching customer data — which is what the FINRA audit trail already is.

What buyers should ask in the procurement process. First: "Show me the unified interaction record format. For one of your reference deployments, show me a single client's six-month interaction history as the platform exports it." Second: "What is the retention guarantee? Does the platform meet 17a-4 WORM standards out of the box, or do we need a separate archiving tier?" Third: "For each of FINRA Rule 3110, FINRA Rule 4511, SEC Rule 17a-4, and SEC Reg BI, name the platform features that support compliance and the customer obligations that remain."

The 2026 FINRA examination programme will continue to focus on AI. Firms with AI deployments are being examined on AI specifically. The firms that survive these examinations are the ones whose audit trails are unified, signed, retained, and ready. The firms that struggle are the ones whose audit trails are still being assembled when the examiner walks in.

PART III

CHAPTER 07

Healthcare

— *HIPAA, FDA SaMD, supervision.*

Three regulatory frames a healthcare buyer has to satisfy simultaneously, and the line between platforms that fit administrative use and platforms that require FDA clearance.

Healthcare is the sector where agentic AI procurement decisions take the longest, and the reason is structural: a healthcare deployment has to satisfy three overlapping regulatory frames simultaneously, none of which were designed with agentic AI in mind. The frames overlap differently for different deployment shapes — an AI assistant that drafts internal emails for hospital administrators sits in a very different regulatory posture from an AI tool that summarises a patient encounter for a clinician. This dossier walks through the three frames, the line that separates productivity tooling from regulated medical devices, and what a healthcare buyer should ask any AI vendor before procurement begins.

Before the regulatory walk-through: a critical scope clarification. Vantage Workspace is a productivity platform with strong audit posture. It is not a clinical decision support system. It is not a medical device under any FDA definition. It does not produce clinical recommendations, does not interpret medical images, does not generate diagnostic conclusions. The platform fits administrative and clinically-adjacent uses (internal communications, document drafting, file management, agent-assisted operational work) where the regulatory exposure is real but bounded. Where a healthcare organisation needs an AI tool that participates in clinical decision-making, that is a different product category — typically requiring FDA 510(k) clearance, ONC certification, or both — and we will tell you so in the first conversation.

The platform fits administrative and clinically-adjacent uses. Where a healthcare organisation needs an AI tool that participates in clinical decision-making, that is a different product category — typically requiring FDA 510(k) clearance.

— SCOPE CLARIFICATION

THREE OPERATIVE FRAMES

HIPAA · FDA SaMD · State medical board

Each applies independently. The platform's posture has to satisfy all three, or name the limit explicitly.

The first frame is HIPAA. The Health Insurance Portability and Accountability Act, with its Privacy Rule (45 CFR Part 164 Subpart E) and Security Rule (Subpart C), is the operative federal framework whenever Protected Health Information is involved. PHI is broadly defined: any individually identifiable health information held or transmitted by a covered entity or business associate, in any form, including electronic. The moment an agentic AI tool reads, drafts, summarises, or transmits PHI, the platform vendor becomes a Business Associate under 45 CFR 160.103, requiring a Business Associate Agreement, the Security Rule's administrative/physical/technical safeguards, breach notification under the HITECH amendments, and the minimum-necessary standard under 164.502(b).

The minimum-necessary standard is the part agentic AI buyers most often underestimate. Under 164.502(b), a covered entity must make reasonable efforts to use, disclose, and request only the minimum amount of PHI necessary to accomplish the intended purpose. Applied to an AI agent, this means: the agent should not have undifferentiated access to a patient's full record when the task at hand only requires a specific section. A chatbot that ingests an entire EHR record to answer a scheduling question fails the minimum-necessary standard, even if every other HIPAA control is in place. The structural answer is per-task scope enforcement at the platform level — which Vantage Workspace's per-agent permission model and policy engine produce, but which is missing from most agentic AI products that ship as "give the agent your data and let it figure out what it needs."

The second HIPAA dimension that catches healthcare deployments: the breach notification calculus when an AI's output is itself the breach. If an agent's reply contains information from another patient's record (a memory leak, an embedding inversion, a retrieval scope error), the disclosure is reportable under 164.404. The notification clock starts at discovery. The HHS Office for Civil Rights has signalled in 2025 enforcement guidance that AI-driven incidents are a focus area. The platform's posture has to assume incidents will happen and produce the contemporaneous evidence record that makes notification fast, accurate, and bounded.

The second frame is FDA Software as a Medical Device. The FDA's authority over software as a medical device derives from the Federal Food, Drug, and Cosmetic Act and is operationalised through the Center for Devices and Radiological Health (CDRH). The relevant guidance for AI/ML-enabled software has matured significantly since the FDA's 2019 discussion paper, with the January 2025 final guidance on Predetermined Change Control Plans (PCCPs) for AI/ML SaMD providing the current procedural framework for how such products are cleared and how they can change post-market without requiring re-clearance.

The line that matters for healthcare AI procurement is whether the software is a medical device under section 201(h) of the FD&C Act. The 21st Century Cures Act narrowed the definition for clinical decision support software at section 3060, exempting CDS that meets four specific criteria. Software that fails any of the four criteria — particularly the "independent review" clause for recommendation-generating tools — falls back into the device definition and requires FDA pathway selection.

Applied to agentic AI, the test is concrete: an agent that drafts a clinical note for a physician to review and edit before signing is administrative; an agent that generates a diagnostic recommendation that the physician relies on without independent verification is potentially device-regulated. The line is sharper than vendor marketing suggests. The 2025 guidance specifically addresses how iterative AI changes (model retraining, prompt template updates, tool catalogue changes) interact with the device pathway: a PCCP allows pre-authorised changes within a declared envelope; changes outside the envelope require new submission.

What this means for procurement: a healthcare organisation evaluating Vantage Workspace for administrative and clinically-adjacent uses faces a HIPAA evaluation but not an FDA evaluation. A healthcare organisation evaluating an agentic AI product for clinical decision support faces both, and should be cautious of vendors who claim FDA exemption without explaining which of the four Cures Act criteria the product satisfies. We do not claim FDA exemption because we do not deploy in clinical decision support; the question does not arise.

The third frame is the supervision and accountability layer. State medical boards have begun publishing guidance on AI-assisted clinical practice, with California's Medical Board (under AB 3030, effective January 2025) requiring physicians to review and confirm AI-generated content before it is communicated to patients, and similar requirements emerging in Texas (HB 4660), New York (the proposed AI in Healthcare Disclosure Act), and the Federation of State Medical Boards' 2024 model guidance. The accountability remains with the licensed clinician; the AI tool is treated as a medical assistant whose work the clinician supervises and signs.

The Joint Commission's standards for hospital accreditation address AI under the Information Management chapter and the Patient Safety Systems chapter, with the 2025 revision adding specific expectations around AI governance, including documented policies for AI tool selection, validation processes for AI outputs in clinical contexts, and incident reporting for AI-related adverse events. Joint Commission surveyors increasingly ask to see the AI inventory and the supervision framework as part of standard surveys, not as a special focus area.

Liability allocation is the question that ultimately determines deployment shape. Under current malpractice doctrine, the supervising physician retains responsibility for clinical decisions even when those decisions are informed or generated by AI. A vendor that fails to deliver the audit trail necessary to demonstrate physician supervision is exposing the physician to a defence-burdening evidence gap. A platform that produces unified, signed, contemporaneous records of AI output and physician review is the platform that supports defensible practice.

The platform's posture for healthcare deployments. Vantage Workspace fits the administrative and clinically-adjacent layer of a healthcare organisation's workload. The platform produces the HIPAA-required evidence record, supports per-agent scope enforcement consistent with the minimum-necessary standard, ships with Keycloak preconfigured (federating to the customer's existing identity provider — typically Okta, Microsoft Entra ID, or Epic-integrated identity in EHR-heavy environments — when one is in place), and feeds the customer's existing SIEM where their security operations are already running.

What the platform does NOT do for healthcare: it does not generate clinical recommendations, interpret diagnostic data, or produce content that a clinician would rely on without independent review. It does not have FDA clearance because it does not require it. It does not replace the EHR. It does not handle the physician supervision workflow on its own — that workflow is configured by the customer's medical informatics team using the platform's policy framework.

Three deployment patterns we see in healthcare procurement. The first is administrative-only: the platform is used for internal hospital operations where PHI is incidentally present but not the focus. HIPAA applies; FDA does not; supervision is general organizational accountability. The second is clinically-adjacent: the platform supports work that touches clinical operations without participating in clinical decisions. HIPAA applies in full; FDA generally does not, provided no clinical recommendations are generated. Most healthcare deployments of agentic productivity platforms land here. The third is clinical-decision-support, where the AI participates in care decisions. This is where Vantage Workspace stops being the right product; we will tell a healthcare buyer this directly.

Four questions a healthcare buyer should ask any agentic AI vendor before procurement. First: "Will you sign a Business Associate Agreement, and which of the Security Rule's safeguards do you implement at the platform level?" Second: "How does the platform enforce the minimum-necessary standard at the per-agent, per-task level? Show me a worked example of an agent's permission scope being narrower than the user who delegated to it." Third: "Is your platform regulated as a medical device under FDA SaMD criteria? If you claim exemption under the Cures Act, which of the four section 3060 criteria does the platform satisfy?" Fourth: "What does the audit record look like for a single physician interaction across a 24-hour shift, and can you produce it during an OCR examination or a Joint Commission survey within the response window the regulator allows?"

Vantage Workspace fits a specific portion of the healthcare AI workload. The portion it fits is large — most administrative and clinically-adjacent productivity work falls in this category. The portion it does not fit (clinical decision support) is a different product category. A healthcare buyer who knows which use cases sit where in this map can make a procurement decision faster, with less risk of either deploying the wrong tool for the job or accepting unnecessary regulatory exposure.

PART III

CHAPTER 08

Fintech

— *BSA/AML, fair lending, sponsor bank.*

Three regulatory pressures fintechs face that incumbents don't, and the line between platforms that fit operations and customer service vs. platforms that need to be the credit-decisioning engine.

Fintech is structurally harder than the financial-services dossier suggests. That earlier piece looked at FINRA-registered broker-dealers and SEC-registered advisers — large, established firms with mature compliance functions and decades-deep regulatory relationships. Fintech is different in three ways that matter for agentic AI procurement: the regulatory map has more pieces and they overlap more confusingly, the supervisory relationship often runs through a sponsor bank rather than directly to a regulator, and the company is typically smaller and faster-moving than the compliance burden alone would suggest is sustainable. This dossier walks through what changes when the fintech CEO, the head of compliance, or the CTO is the buyer.

Before the regulatory walk: a critical scope clarification, parallel to the one in our healthcare dossier. Vantage Workspace is a productivity platform with strong audit posture. It is not a credit-decisioning system. It does not produce underwriting recommendations, does not perform automated KYC adjudication, does not generate fraud-decision outputs that go directly to a customer without human review. The platform fits internal operations, employee-facing AI workflows, and customer-service-support use cases. It does not fit the use case where the agentic AI itself is the decisioning engine in a credit, lending, payments, or AML workflow. We will tell a fintech buyer this directly in the first conversation.

A fintech that cannot demonstrate a documented AI governance posture, with audit trails the bank's examiners can review, is increasingly being deplatformed or denied new partnerships.

— THE POST-SYNAPSE SPONSOR-BANK TIGHTENING

INTERAGENCY GUIDANCE

OCC · FDIC · Fed · June 2023

Joint third-party risk management framework. Sponsor banks now apply it to fintechs as a matter of course — and to the AI tools fintechs use.

The first regulatory frame is BSA and AML. The Bank Secrecy Act and the related anti-money-laundering regime are the operative federal framework for any fintech moving customer funds, opening customer accounts, or processing payments. FinCEN administers BSA; specific obligations include the Customer Identification Program (31 CFR 1020.220), the Customer Due Diligence rule (1010.230), Suspicious Activity Reporting (1020.320), Currency Transaction Reports, and the OFAC sanctions screening regime. Each of these has implications for how AI tools can be used in customer onboarding, transaction monitoring, and risk scoring.

The FinCEN guidance from 2024 on AI in BSA compliance programs makes the operative principle explicit: AI tools can augment a BSA compliance program but cannot be the sole decisioning layer for SAR filings, for sanctions hits, or for high-risk customer designations. Human review at the moment of decision is required, and the audit trail has to demonstrate that human review actually happened — not just that a human was theoretically in the loop. The platform that supports this requirement is the platform that produces a unified record of the AI's output, the human's review timestamp, the human's reasoning, and the final decision.

The second regulatory frame is fair lending. The Equal Credit Opportunity Act (ECOA) and its implementing Regulation B (12 CFR 1002) prohibit discrimination on prohibited bases in any aspect of a credit transaction. The Fair Housing Act extends similar protections to mortgage lending. The Consumer Financial Protection Bureau is the operative enforcer for non-bank lenders; bank lenders face overlapping enforcement from their primary federal regulator and the CFPB.

The CFPB's September 2023 statement on adverse action notices for AI-driven credit decisions changed the operating burden materially. Under Regulation B and the Fair Credit Reporting Act, lenders must provide adverse action notices that include specific principal reasons for the denial. The CFPB clarified that generic reason codes are not acceptable when an AI model used dozens or hundreds of features in the decision. The notice must reflect the actual factors that drove the model's decision for that specific applicant. This is operationally hard for opaque models and impossible for models without proper output-to-feature attribution.

Where Vantage Workspace fits in the fair-lending picture: the platform supports fintech employees in operational work that is adjacent to lending (drafting customer communications, summarising compliance reports, coordinating with the compliance team) but does not generate the credit decision itself. The audit trail the platform produces is what supports the fintech's ability to demonstrate to an examiner that AI was used for permitted purposes and not for prohibited ones. Where a fintech needs an AI tool that participates in the credit decision, that tool needs to come from a vendor specialising in explainable AI underwriting (companies like Zest AI, Upstart's licensable model platform, Stratyfy, or LenddoEFL).

The third regulatory frame is the sponsor-bank relationship for fintechs operating under a Banking-as-a-Service (BaaS) or middleware model. A non-bank fintech offering bank-like services typically does so through a partnership with a chartered bank that provides the regulatory umbrella. The sponsor bank's primary federal regulator (OCC for national banks, FDIC for most state-chartered banks, the Federal Reserve for state member banks) treats the fintech as a third party to the bank — and applies the full third-party risk management framework to it.

The Interagency Guidance on Third-Party Relationships, jointly issued by the OCC, FDIC, and Federal Reserve in June 2023 (replacing earlier OCC Bulletin 2013-29 and similar predecessors), is the operative document. It requires the sponsor bank to perform due diligence on the fintech's information security, business continuity, compliance program, and operational resilience — including how the fintech uses AI tools. After the high-profile failures of 2024 (the Synapse collapse, the Evolve Bank consent order, multiple OCC enforcement actions against sponsor banks for inadequate fintech oversight), sponsor banks have tightened their due diligence requirements significantly. A fintech that cannot demonstrate a documented AI governance posture is increasingly being deplatformed or denied new partnerships.

This is where Vantage Workspace addresses a specific pain point that incumbents don't face. A large bank already has a model risk management program, an information security organisation, and decades of audit infrastructure. A fintech with 40 employees, an aggressive product roadmap, and a sponsor bank asking for SR 11-7-equivalent documentation is structurally underserved by the productivity tooling available to it. The platform's value proposition for this fintech is the documented AI governance posture, the unified audit trail, and the framework-mapped compliance evidence — produced as a feature of the platform, not as a quarterly project the fintech's small compliance team has to construct from scratch.

Three deployment patterns we see in fintech procurement. The first is operations-only: the platform is used for internal company operations where customer data is incidentally present but not the focus. The second is customer-service-support: the platform supports employees who interact with customers, helping them draft responses, summarise account history, coordinate across teams. The AI does not make decisions the customer sees directly — the employee reviews and sends. UDAAP applies (because the platform's outputs reach customers via the employee). Most fintech deployments of agentic productivity platforms ultimately land here. The third is decision-making AI in a regulated workflow — credit decisioning, AML transaction monitoring, KYC adjudication, fraud-decision automation. This is where Vantage Workspace stops being the right product.

What the customer still owns. The BSA compliance program is the fintech's; we provide the platform-level audit infrastructure, but the SAR review workflow, the high-risk customer determination, the OFAC hit adjudication remain the compliance team's work. The fair-lending policy and the underwriting guidelines (where the fintech is doing any kind of lending) are the fintech's. The relationship with the sponsor bank — including the documentation the bank requires for AI governance — is the fintech's; we make the documentation easier to produce, but we do not deliver the bank relationship.

Five questions a fintech buyer should ask any agentic AI vendor before procurement. First: "Will you sign a data processing agreement that satisfies our sponsor bank's third-party risk management requirements? Specifically, will you commit to the controls in the OCC's 2023 Interagency Guidance?" Second: "For BSA/AML use cases, how does the platform support the requirement that AI tools augment but do not replace human decisioning?" Third: "Does any output of your platform reach a customer without human review? If so, walk me through how that output is monitored for UDAAP risk." Fourth: "What is your model risk management documentation? Specifically, can you provide what an OCC examiner would want to see under SR 11-7 extended to AI/ML models?" Fifth: "What does the audit record look like during a CFPB examination?"

The 2026 fintech regulatory environment is being shaped by three converging pressures: the post-Synapse tightening of sponsor-bank due diligence, the CFPB's expanded focus on algorithmic accountability in consumer financial products, and the OCC's extension of SR 11-7 model risk management expectations to AI/ML across the regulated banking system. Fintechs that will continue to grow in this environment are the ones that can demonstrate documented AI governance to their sponsor bank, can produce examination-ready audit trails to their consumer regulator, and can match the right AI tool to the right use case without overreaching into regulatory categories the tool was not designed for.

Vantage Workspace fits a specific portion of the fintech AI workload: operations, internal coordination, and customer-service-support — the work that touches the regulatory perimeter without being the regulated decision itself. The portion is large enough that most fintechs we talk to could deploy the platform across 60–80% of their AI use cases. The portion the platform does not fit is a different product category, and a fintech buyer who knows the difference can make procurement decisions that survive sponsor-bank scrutiny and consumer-regulator examination.

PART III

CHAPTER 09

Canadian public sector — *TBSDADM,* *sovereignty, Indigenous* *data.*

Five overlapping jurisdictional frames a Canadian public-sector buyer satisfies simultaneously, why "data residency" isn't a one-line answer for First Nations governments, and what platform sovereignty actually means in this context.

Canadian public-sector AI procurement is structurally different from the US federal model the agentic-AI vendor ecosystem has spent the last two years optimising for. The differences matter at the procurement-decision level: a federal department in Ottawa, a provincial agency in Quebec City, a Crown Corporation operating across the country, and a First Nations government with a self-government agreement are all considered "public sector" in casual conversation, but they answer to different regulators, operate under different privacy frameworks, hold different procurement obligations, and — in the case of First Nations governments — exist within a sovereignty frame that requires different conversations entirely.

Before the regulatory walk: a critical scope clarification. Vantage Workspace is a productivity platform for internal government operations and employee-facing AI workflows. It is not a citizen-facing decision-making system. It does not produce automated decisions that bind individuals (a benefits eligibility determination, a permit denial, an immigration adjudication). The federal Treasury Board Directive on Automated Decision-Making and provincial equivalents apply specifically to systems that make or recommend decisions affecting the rights, privileges, or interests of individuals. The platform fits internal coordination, document drafting, briefing-note preparation, and operational work; it does not fit the use case where the AI itself is the decisioning layer in a citizen-facing service.

Vendors who treat Indigenous data sovereignty as a checklist item will lose deals, often quickly and quietly. Vendors who approach the conversation with humility about the limits of what generic compliance frameworks can address will earn the conversation.

— ON FIRST NATIONS DATA SOVEREIGNTY

FIVE JURISDICTIONAL FRAMES

Federal · Provincial · Crown Corp · Sovereign Nations · Sensitive institutions

Each operates under different rules. A vendor that treats Canada as one jurisdiction will struggle in any of the five.

The first regulatory frame is federal. The Treasury Board Secretariat's Directive on Automated Decision-Making (TBSDADM, in force since April 2019, amended April 2023) is the operative federal framework for any automated system that supports or replaces an administrative decision affecting an individual. The Directive establishes four impact-level classifications based on the reversibility, duration, and reach of the decision's effects, with progressively stringent requirements at each level — including the mandatory Algorithmic Impact Assessment (AIA) before deployment. The 2023 amendments tightened the requirements for transparency, peer review, and ongoing monitoring.

The federal privacy frame layers on top. The Privacy Act (R.S.C. 1985, c. P-21) governs the collection, use, and disclosure of personal information by federal institutions; the Personal Information Protection and Electronic Documents Act (PIPEDA) governs commercial-context personal information. The Office of the Privacy Commissioner of Canada has issued multiple AI-specific guidance documents since 2023. AIDA — the Artificial Intelligence and Data Act, currently moving through Parliament — would extend formal AI obligations to private-sector deployments, but federal institutions will continue to be governed primarily by the Privacy Act and the TBSDADM regardless of AIDA's final form.

The second frame is provincial. Each province has its own privacy framework, and the patchwork is real. Quebec's Law 25 (formerly Bill 64, fully in force since September 2024) is the most stringent, with explicit AI provisions including a transparency requirement when an automated decision is made about an individual and a right of human review. British Columbia and Alberta operate under PIPA. Ontario's public sector operates under FIPPA and MFIPPA; the Ontario AI Framework, established by directive in December 2023, lays out the operational expectations. The Information and Privacy Commissioners of Ontario, BC, and Quebec each have published AI-specific guidance and an examination posture, with growing focus on AI in 2025 and 2026.

The third frame is the Crown Corporation dimension. Crown Corps occupy a particular space: they are federally chartered (or provincially, in some cases) but operate with commercial-style mandates and degrees of independence that vary by enabling legislation. Most Crown Corps are subject to the Privacy Act and (where the Directive applies) the TBSDADM, but the application is typically through their enabling legislation rather than directly. AI procurement at a Crown Corp involves the same federal frameworks as a department but with the addition of the Corp's own commercial considerations.

The fourth frame, and the one most non-Indigenous procurement teams underestimate, is sovereign Nations. There are over 600 First Nations communities in Canada, each with its own government structure; some operate under the Indian Act, others have negotiated self-government agreements that establish jurisdictional authority over their own data, services, and operations. Métis nations and Inuit governance bodies have parallel but distinct sovereignty frames. The Federal Court of Appeal's 2024 jurisprudence on First Nations data sovereignty, the OCAP® principles articulated by the First Nations Information Governance Centre (Ownership, Control, Access, and Possession), and the United Nations Declaration on the Rights of Indigenous Peoples (UNDRIP, adopted into Canadian law via the UNDRIP Act in 2021) collectively establish that data about Indigenous individuals and communities is governed by frameworks that are not reducible to PIPEDA or provincial privacy law.

Practically, what this means for an AI vendor: a procurement conversation with a First Nations government, a Crown corporation operating in Indigenous-relations contexts, or a federal program serving Indigenous communities is not a procurement conversation that "the platform satisfies our standard data residency requirements" will close. The conversation requires consultation with the Nation's data governance body (where one exists), respect for OCAP principles in the platform's deployment, and — in the case of self-governing Nations — recognition that the Nation itself is the regulator of how data about it and its members is processed. Vendors who treat Indigenous data sovereignty as a checklist item will lose deals, often quickly and quietly.

The fifth frame is core sensitive institutions. This category — defence-adjacent organisations (Department of National Defence, Communications Security Establishment, Canadian Security Intelligence Service), critical infrastructure operators, public safety institutions, and certain health system bodies — operates under additional security frameworks that overlay the standard public-sector regime. The Government of Canada Security Categorization Standard, the Communications Security Establishment's Cyber Security Centre guidance, and (for defence-specific contexts) the requirements applicable to controlled goods and ITAR-equivalent regimes establish what AI deployments are permitted and what platform-level controls are required.

The platform's posture for Canadian public sector. Vantage Workspace fits the internal-operations and employee-facing AI workflow layer of a public-sector institution's workload. The platform produces the audit evidence record that satisfies AIA monitoring requirements where applicable, supports per-agent scope enforcement consistent with the minimum-collection principles in federal and provincial privacy law, ships with Keycloak preconfigured (federating to the institution's existing identity provider — typically Active Directory / Microsoft Entra ID for federal departments, GCKey or provincial equivalents for citizen-facing systems — when one is in place), and feeds the institution's existing SIEM where their security operations are running.

What the platform does NOT do for Canadian public sector: it does not produce automated decisions that fall under the TBSDADM's Level III or Level IV impact classifications. It does not handle the Algorithmic Impact Assessment process on the institution's behalf — the AIA is the institution's document. It does not replace ProtectedB or higher security classification handling — the platform's deployment must align with the security categorization of the data it will touch. For Indigenous data contexts specifically, the platform deployment must be negotiated with the relevant Nation or governance body; we do not claim a generic OCAP compliance posture, because OCAP compliance is a relationship, not a vendor checklist.

Three deployment patterns we see. The first is internal-operations-only: briefing note drafting, internal coordination, vendor management, project documentation. The TBSDADM does not apply (no automated decisions affecting individuals); the Privacy Act and provincial equivalents apply where personal information is incidentally present. The second is employee-facing-with-citizen-context: the platform supports employees who interact with citizens or who handle citizen data. The AI does not make decisions citizens see directly. Most Canadian public-sector deployments eventually land here. The third is automated-decision systems falling under the TBSDADM's Level III or IV classifications. Vantage Workspace stops being the right product here.

Five questions a Canadian public-sector buyer should ask any agentic AI vendor before procurement. First: "What is your security clearance status, and at what classification levels can you support deployment? Specifically, can you support a deployment at ProtectedB?" Second: "For the TBSDADM, where in your platform do you produce technical evidence that supports the Algorithmic Impact Assessment process?" Third: "For provincial deployments, particularly in Quebec under Law 25, how does your platform support the transparency and human-review obligations?" Fourth: "Have you deployed in or alongside a First Nations government, Crown Corporation, or Indigenous-relations context? If so, how did you approach the OCAP principles? If not, are you willing to work through that conversation rather than treat it as standard compliance?" Fifth: "What does the audit record look like for an Office of the Information and Privacy Commissioner examination, or for a Treasury Board Secretariat audit?"

Vantage Workspace fits a specific portion of the Canadian public-sector AI workload. The portion is large — most internal-operations and employee-facing work in a department, agency, Crown Corporation, or Indigenous government falls in this category — and the regulatory burden is real but bounded. The portion it does not fit is a different product category, and a public-sector buyer who knows the difference can make procurement decisions that survive Office of the Privacy Commissioner examination, Treasury Board Secretariat audit, and the consultation obligations the Crown and individual institutions hold to Indigenous Nations and communities.

PART III

CHAPTER 10

Legal services — *privilege, competence, supervision.*

Three professional-conduct rules that change when AI processes privileged communications, what ABA Formal Opinion 512 and the Canadian law societies actually require, and the line between AI as a research tool and AI as the practice of law.

Legal services is the sector where agentic AI procurement has the smallest margin for error, because the same conduct rules that make a lawyer's work valuable — confidentiality, competence, supervision, the duty to communicate honestly with clients — apply unmodified to AI tools that participate in the work. A misjudged AI deployment in a law firm is not a productivity issue; it is a discipline-committee issue, a malpractice-insurance issue, and a reputational issue all at once. Three lawyers were sanctioned in the Avianca case for ChatGPT-fabricated citations in 2023, and the discipline cases have continued through 2024 and 2025 in growing volume.

Before the regulatory walk: a critical scope clarification, parallel to the healthcare and fintech dossiers. Vantage Workspace is a productivity platform that supports lawyers in their work. It is not a substitute for a lawyer. It does not produce legal advice that a client receives directly. It does not draft and file documents that go to court without lawyer review. It does not engage in the practice of law, which in every jurisdiction is reserved for licensed members of the bar. The platform fits the workflow of a lawyer or a legal professional doing their work: drafting, research, summarisation, scheduling, document management, internal coordination. The platform does not fit the unauthorised-practice-of-law concern that has driven discipline against several AI-assisted "legal service" companies.

The conduct rules do not prohibit AI assistance — they require that the lawyer's supervision actually happen, with documentation that demonstrates it happened, before the work product leaves the firm.

— ON MATA V. AVIANCA AND THE SUPERVISION RULE

ABA FORMAL OPINION

512 • July 2024

Consolidates confidentiality (1.6), competence (1.1), supervision (5.3), and client communication into one operative framework for generative AI.

The first conduct rule is confidentiality. The American Bar Association's Model Rule 1.6 and the parallel rules in every state of the US (and the equivalent provisions in the Federation of Law Societies of Canada Model Code, adopted with provincial variations) impose a duty on lawyers to preserve the confidentiality of information relating to the representation of a client. The rule extends beyond solicitor-client privilege to all information about the representation, regardless of whether the information is protected by privilege rules in court.

When an agentic AI tool processes confidential information, the question becomes: where does that information go, and who can see it? Most consumer-grade AI assistants and most enterprise SaaS AI products process data on the vendor's infrastructure, with terms of service that permit some level of telemetry, training-data use, or vendor employee access for support purposes. None of those terms are necessarily incompatible with Rule 1.6, but they require analysis: the lawyer needs to understand where the data goes, what the vendor's contractual restrictions are, whether informed client consent is required, and whether the lawyer's obligation to take reasonable precautions to prevent inadvertent disclosure (Comment [18] to Rule 1.6 in the US, similar comments in Canadian rules) is satisfied by the vendor's security posture.

The Law Society of Ontario's 2024 guidance on AI use in legal practice, the Canadian Bar Association's 2024 practice resource, and the Federation of Law Societies of Canada's 2025 model code amendments collectively make the operative principle explicit: the lawyer remains responsible for the confidentiality of information processed by AI tools the lawyer chooses to use. Vendor terms that permit any use of client data outside the immediate processing of the lawyer's request — for model training, for product improvement, for vendor analytics — are likely to fail Rule 1.6 unless the client has been informed and has given informed consent.

The second conduct rule is competence. ABA Model Rule 1.1 (and the parallel competence rules in Canadian provincial codes) requires lawyers to provide competent representation. Comment [8] to Rule 1.1, added in 2012 and now adopted in nearly every US jurisdiction, expressly extends competence to include keeping abreast of changes in the law and its practice — "including the benefits and risks associated with relevant technology."

Appplied to agentic AI, the competence rule means a lawyer who chooses to use an AI tool has an affirmative duty to understand what the tool does, where its limitations lie, and what risks attach to its use. The ABA Formal Opinion 512, issued in July 2024, made this explicit for generative AI specifically: lawyers must "reasonably understand" the AI tools they use, including whether the tool may produce false or misleading outputs (the "hallucination" problem), whether the tool retains or transmits client information, and whether the tool's use is consistent with the lawyer's other professional obligations. A lawyer who delegates work to an AI tool without this understanding is at risk of a competence violation regardless of whether the AI's output is correct.

The third conduct rule is supervision. ABA Model Rule 5.1 governs the supervision of subordinate lawyers; Rule 5.3 governs the supervision of non-lawyer assistants. The 2024 amendments to the comments on Rule 5.3 (and parallel guidance in Canadian provinces) clarified that AI tools used in legal practice are governed by the supervision framework that applies to non-lawyer assistants. The lawyer is responsible for the AI's work as if the work had been done by a paralegal or assistant under the lawyer's supervision.

The supervisory framework in practice means: the lawyer is responsible for verifying the AI's outputs before they are used. An AI that drafts a research memo with fabricated case citations has, under the supervision rule, produced work the supervising lawyer had a duty to verify. The *Avianca* decision (*Mata v. Avianca Inc.*, S.D.N.Y. 2023) and the discipline cases that followed it each turned on this point: the lawyers had not verified the AI's output before submitting it to the court. The conduct rules do not prohibit AI assistance — they require that the lawyer's supervision actually happen, with documentation that demonstrates it happened, before the work product leaves the firm.

The privilege dimension is distinct from confidentiality and worth naming separately. Solicitor-client privilege (in Canadian law) and attorney-client privilege (in US law) protect communications between a lawyer and client made for the purpose of obtaining legal advice, from compelled disclosure in court. The privilege belongs to the client, not the lawyer. The current consensus across US and Canadian jurisdictions is that processing through a vendor's AI tool does NOT automatically waive privilege provided the vendor is under a confidentiality obligation equivalent to the lawyer's, and provided the vendor's access is limited to what is necessary to perform the contracted service.

However, the privilege analysis becomes more complicated when (a) the AI vendor uses the data for purposes beyond the immediate processing, (b) the AI vendor's subcontractors or sub-processors have access, or (c) the AI vendor's jurisdiction of operation creates exposure to compelled disclosure under foreign law (a particular concern with US-based vendors processing Canadian-client data in a context where US discovery orders might reach the vendor). The Canadian Bar Association's 2024 practice resource specifically flags the cross-border data processing concern as a privilege risk that lawyers should evaluate before deploying AI tools.

ABA Formal Opinion 512 (July 2024) is the most significant recent guidance and consolidates the analysis across confidentiality, competence, supervision, and the duty of communication with clients. The opinion specifically addresses generative AI use and concludes: lawyers may use AI tools, must understand them, must take reasonable precautions to protect client information, must verify the AI's outputs, must obtain client consent in some circumstances, and must consider whether to disclose AI use in the lawyer's billing and communication with the client. Subsequent state ethics opinions in California, New York, Florida, Illinois, and several other jurisdictions have largely tracked Opinion 512's approach.

The platform's posture for legal services. Vantage Workspace fits the workflow of lawyers and legal professionals doing their work, with a deployment model that addresses each of the four professional-conduct dimensions. On confidentiality (Rule 1.6): the platform runs on the firm's infrastructure, the firm controls the data plane, no telemetry phones home to the vendor, and no data is used for model training. On competence (Rule 1.1): the platform's outputs are accompanied by sufficient transparency for the lawyer to verify (model version, prompt, retrieval sources where applicable). On supervision (Rule 5.3): the platform's audit log records the lawyer's review action — when the lawyer reviewed the AI's output, what changes the lawyer made, when the lawyer signed off. On privilege: the platform's deployment on firm infrastructure means no third-party AI vendor accesses client data in the way that would create privilege risk.

What the platform does NOT do for legal services: it does not provide legal advice. It does not file documents on a client's behalf. It does not engage with clients without lawyer involvement. It does not produce work product that a court receives without lawyer review and signature. It does not adjudicate matters. Each of these would constitute the practice of law, which is reserved to licensed lawyers; the platform is designed to support lawyers, not to replace them.

Three deployment patterns we see in legal-services procurement. The first is internal-firm-operations: knowledge management, internal communications, conflicts checking, scheduling, billing administration. Confidentiality applies in the standard way; privilege is not implicated because the work is not legal practice. The second is lawyer-augmenting-on-client-matters: research, drafting, document review, summarisation of case materials, where the lawyer's judgment and supervision are central. Most legal-services deployments of agentic productivity platforms ultimately land here. The third is client-facing-AI or AI-as-the-practice-of-law. This is where Vantage Workspace stops being the right product.

Five questions a legal-services buyer should ask any agentic AI vendor before procurement. First: "Where is client data processed, and what are the contractual restrictions on the vendor's use of that data? Specifically, can you confirm in writing that no client data is used for model training?" Second: "For ABA Model Rule 1.1 competence, what materials do you provide to help our lawyers reasonably understand your AI tool?" Third: "For Rule 5.3 supervision, what does the audit log look like for a lawyer's review of an AI output? Walk me through how a discipline-committee investigation would reconstruct, six months after the fact, that a specific lawyer verified a specific AI output before it was used." Fourth: "What is your sub-processor list, where do they operate, and what is your exposure to compelled disclosure under foreign law?" Fifth: "Have you been the subject of any discipline-committee investigation, malpractice-insurance claim, or court sanction related to AI use in legal practice?"

Vantage Workspace fits a specific portion of the legal-services AI workload — internal firm operations and lawyer-augmenting work where the lawyer's judgment remains central. The portion is large enough that most law firms could deploy across most of their internal AI use cases. The portion the platform does not fit (client-facing legal-advice automation, document filing without lawyer review) is a different product category, with its own regulatory pressure and its own discipline risks. A legal-services buyer who knows the difference can make procurement decisions that survive ethics review, malpractice underwriting, and the discipline-committee scrutiny that will continue to define the AI-in-legal-practice conversation through 2026 and beyond.